



**L**OVELY  
**P**ROFESSIONAL  
**U**NIVERSITY

## **SUMMER TRAINING/INTERNSHIP**

### **PROJECT REPORT**

**(Term June-July 2025)**

**Title of Project : (E-commerce Product Recommendation System.)**

**Submitted by**

**Name of Student : Kasala Hoshik.**

**Registration Number : 12306410**

**Name of Student : Badana Lohith.**

**Registration Number : 12318788**

**Name of Student : Atmakuri Rama Venkata Akash Gupta.**

**Registration Number : 12310650**

**Name of Student : Uppumaguluru Aravind.**

**Registration Number : 12322664**

**Name of Student : Y.Mahendra.**

**Registration Number : 12317546**

**Course Code : PETV79 – Machine Learning Made Easy :  
From Basics to Ai Applications.**

**Under the Guidance of**

**Name of mentor : Mahipal Singh Papola.**

**UID : 32137**

**School of Computer Science and Engineering**

# CHAPTER 1: INTRODUCTION

## COMPANY PROFILE

The internship for this project was undertaken under the flagship training initiative titled “Machine Learning Made Easy: From Basics to AI Applications.” This program was designed and hosted by Lovely Professional University, committed to bridging the gap between academic learning and real-world implementation in the field of Artificial Intelligence and Machine Learning.

The organization behind this initiative operates with a vision to make AI education approachable and practical. By offering project-oriented internships, the program gives budding developers, engineers, and analysts the opportunity to work on real datasets and apply algorithms to solve impactful problems. Over the years, the organization has built a reputation for mentoring thousands of students across the country and encouraging innovation in emerging technologies like machine learning, data science, and deep learning. The environment fostered by this training body emphasizes collaborative learning, guided experimentation, and end-to-end project execution.

The program welcomes students from a variety of technical backgrounds, helping them gain proficiency in key ML concepts while also strengthening their confidence in using libraries, frameworks, and cloud-based tools. With a strong focus on project-based assessments, participants walk away not just with certificates, but with deployable portfolios and deeper insight into industry standards.

## OVERVIEW OF TRAINING DOMAIN

The core focus of our internship revolved around the domain of Recommendation Systems, a subfield of machine learning that has gained immense traction due to its use in modern-day applications ranging from Netflix and Spotify to Amazon and Flipkart.

Recommendation systems are designed to offer personalized suggestions to users based on their previous activity, product characteristics, or the behavior of similar users. Broadly, the field is divided into three categories:

- **Collaborative Filtering (CF):** Uses past user behavior and ratings to find patterns and recommend new items.
- **Content-Based Filtering (CB):** Focuses on the properties of items and user preferences to find similar items.
- **Hybrid Systems:** Combine both CF and CB to deliver more accurate and reliable recommendations.

During our training, we were exposed to all three paradigms. The learning was progressive, starting from the basics of Python programming and data handling, and moving into deeper territory like TF-IDF vectorization, cosine similarity, matrix factorization, model evaluation metrics, and real-time data parsing.

What made this domain especially exciting was its clear relevance to our everyday digital experience. The importance of recommending the right product to the right person at the right time is a core challenge for many businesses today, especially those in the e-commerce and content delivery space. It was also fascinating to learn about the cold start problem, sparsity challenges, and the trade-offs between interpretability and performance in recommender algorithms.

As part of this domain, we learned not just how to build models, but how to design an entire recommendation pipeline — including data preprocessing, feature extraction, training/testing split, model selection, evaluation, and presenting results in a user-friendly format. This gave us an end-to-end view of how data science solutions are deployed in the industry.

## **OBJECTIVE OF THE PROJECT**

The objective of our project was to build a robust and intelligent e-commerce product recommendation system using both collaborative and content-based filtering techniques. The system should be capable of providing accurate and meaningful product suggestions to users based on either their past behavior or textual search input.

This project was designed to simulate a real-world e-commerce experience and address a common business need: how can we help users discover the most relevant products quickly and efficiently?

### **The goals we defined for the project included:**

- Cleaning and integrating multiple raw datasets from the Olist platform
- Developing collaborative filtering models using the Surprise library (SVD, NMF, KNN Basic)
- Extracting product metadata using natural language processing (TF-IDF)
- Creating a hybrid engine that switches intelligently between collaborative and content-based filtering
- Handling cold-start problems for new users
- Designing an interface where users can enter their ID and get top-N recommendations

Throughout this journey, we learned how to approach machine learning not just as an algorithmic exercise, but as a holistic system design challenge. We dealt with real data problems, ran performance benchmarks, compared model accuracies, and fine-tuned the experience based on what worked best for different types of users.

By the end of the internship, the goal was not just to demonstrate technical ability, but to develop a fully working prototype of a personalized recommendation engine that could be realistically used in a digital commerce platform.

---

## CHAPTER 2: TRAINING OVERVIEW

### Introduction to Training

This chapter outlines the structure, content, and learning trajectory of our internship training program, “Machine Learning Made Easy: From Basics to AI Applications.” The training was crafted to build a strong foundational understanding of machine learning concepts and to enable participants to apply these concepts practically through projects. Each phase of the training was strategically planned to ensure a smooth transition from theoretical understanding to practical application, leading up to the implementation of a real-world recommender system.

Our training sessions were structured with an emphasis on problem-solving, active coding, and peer collaboration. The sessions were delivered through a blend of online webinars, coding labs, and weekly review meetings with our assigned mentor. Daily tasks and goals were set, and progress was monitored through version control using GitHub. This systematic approach made the learning process highly organized and result-oriented.

### Tools & Technologies Used

The tools and technologies we learned and used during our training played a crucial role in the success of our project. These tools provided us with a robust infrastructure for data handling, visualization, model development, and evaluation.

1. Python: We used Python as our primary programming language due to its simplicity, vast libraries, and community support.
2. Google Colab: Enabled us to work in a cloud-based Jupyter environment that supports real-time collaboration and has access to GPUs for model training.
3. Pandas & NumPy: Essential for data manipulation, transformation, and numerical computations.
4. Matplotlib & Seaborn: Helped visualize data distributions, patterns, and performance metrics.
5. scikit-learn: Provided the necessary tools for preprocessing data and implementing machine learning algorithms such as TF-IDF and cosine similarity.
6. Surprise Library: A powerful Python library specifically tailored for building and analyzing recommender systems. We used it for collaborative filtering techniques like SVD and NMF.

### Areas Covered During Training

The training program was comprehensive and multifaceted. Below is an overview of the key areas covered:

- Data Preprocessing: We learned how to clean, merge, and preprocess large datasets. This included dealing with missing values, removing duplicates, handling categorical variables, and converting data types.

- Exploratory Data Analysis (EDA): Visualizing data using statistical plots to understand underlying patterns, trends, and anomalies.
- Machine Learning Fundamentals: We revisited concepts like supervised vs. unsupervised learning, overfitting vs. underfitting, and model evaluation techniques.
- Text Processing: TF-IDF vectorization and tokenization helped us convert unstructured text data into a format suitable for machine learning models.
- Collaborative Filtering Techniques: We explored model-based CF algorithms such as SVD (Singular Value Decomposition), NMF (Non-negative Matrix Factorization), and neighborhood-based methods like KNN.
- Content-Based Filtering: Focused on extracting relevant product features and computing cosine similarity between feature vectors.
- Hybrid Recommendation Strategies: We combined both collaborative and content-based approaches to address limitations like cold-start and data sparsity.
- Performance Evaluation: Learned to use RMSE, MAE, precision, and recall to evaluate and compare model performance.

## Daily Work Summary

The following is a breakdown of our progress over the 6-week internship period:

**Day 1:** - Introduction to recommender systems and the project overview. - Familiarization with the Olist dataset. - Basic data cleaning: removing NaNs, duplicates, typecasting. - Initial exploration of customer, order, product, and review datasets.

**Day 2:** - Feature extraction and joining datasets. - Data visualization: bar plots, histograms, scatter plots, heatmaps. - Learned about and experimented with TF-IDF vectorizer and cosine similarity. - Started implementing the content-based recommendation engine.

**Day 3:** - Introduced to collaborative filtering and the Surprise library. - Trained SVD, NMF, and KNNBasic models. - Evaluated models using RMSE and MAE. - Compared baseline models to model-based CF algorithms.

**Day 4:** - Designed the hybrid recommendation logic. - Combined CF and CB recommendations for cold-start handling. - Implemented logic to dynamically switch strategies based on user history. - Started building the user interface in Colab with inputs for user ID and keywords.

**Day 5:** - Performed parameter tuning (e.g., n-gram ranges, min\_df, max\_df). - Created visualizations to demonstrate model results. - Tested the system with multiple user IDs and search terms. - Identified limitations and documented observations.

**Day 6:** - Conducted final testing of the hybrid system. - Compiled results, screenshots, and model performance charts. - Wrote documentation and prepared for project presentation. - Peer-reviewed code and feedback sessions with mentor.

What We Learned

This phase of training was instrumental in developing both technical and soft skills. From a technical perspective, we mastered the end-to-end lifecycle of a machine learning project — from data preparation to deployment-ready recommendation output. On the soft skills side, we improved our team coordination, documentation skills, presentation abilities, and time management.

Every week came with new challenges, from bugs in code to understanding complex model outputs. However, these difficulties shaped our learning. For instance, while tuning the collaborative filtering models, we encountered high RMSE values due to sparse matrices. Solving this by increasing the number of epochs and normalizing ratings taught us about model behavior and sensitivity.

By the end of our training, we were comfortable with working in a professional-like environment, handling large datasets, experimenting with models, and critically evaluating outputs. This experience has paved the way for us to pursue future projects and internships in machine learning with confidence and clarity.

## CHAPTER 3: PROJECT DETAILS

### Project Title

E-Commerce Product Recommendation System: A Hybrid Approach

### Problem Statement

In today's e-commerce landscape, customers are overwhelmed with the vast number of products available online. Without an intelligent recommendation engine, users struggle to find relevant items quickly, often leading to poor engagement and missed sales opportunities for businesses. The absence of personalization results in a generic browsing experience, where all users are exposed to the same set of featured products, irrespective of their unique interests or purchase history.

This creates a dual challenge: first, to design a model that effectively captures and analyzes user behavior, and second, to align it with the metadata of products for users with limited interaction history. Our project tackles this issue head-on by building a robust hybrid recommendation system that combines collaborative filtering with content-based filtering. This system ensures accurate, dynamic, and personalized product suggestions for each individual user.

### Scope of the Project

The project focuses on designing, building, and evaluating a machine learning-based product recommendation system using real-world data from the Olist e-commerce platform. It involves data cleaning, integration of multiple datasets, user-product interaction analysis, and recommendation generation using advanced techniques.

### The scope includes:

- Developing a recommendation engine that can work for both high-activity and cold-start users
- Analyzing user ratings and feedback to inform suggestions
- Extracting semantic insights from product metadata to enable similarity-based recommendations
- Evaluating models using industry-standard metrics
- Providing a framework that is scalable and can be extended to real-time systems

### Objectives of the Project

1. Data Understanding and Preprocessing
  - Merge and clean datasets such as customers, orders, products, reviews, and payments.
  - Handle missing values, duplicates, and inconsistencies.



2. Collaborative Filtering Implementation
  - Train models like SVD, NMF, KNNBasic using the Surprise library.
  - Evaluate performance using RMSE and MAE.
3. Content-Based Filtering Module
  - Extract product features using TF-IDF from product names, brands, tags, and descriptions.
  - Compute cosine similarity to find similar items.
4. Hybrid Model Development
  - Integrate both filtering techniques.
  - Build logic to dynamically switch based on user profile and history.
5. Testing & Evaluation
  - Test model predictions on real users.
  - Generate performance visualizations and accuracy metrics.
6. User Interaction Layer
  - Enable keyword-based product search.
  - Provide top N personalized recommendations.

## System Architecture

The system architecture is composed of the following layers:

- **Data Layer:** Responsible for ingesting raw datasets from Olist. These include customers, orders, reviews, products, and category metadata.
- **Preprocessing Layer:** Cleans, merges, and transforms data into model-ready formats. This includes text normalization, numerical feature engineering, and matrix construction.
- **Modeling Layer:** Houses collaborative filtering models and content-based similarity engines. This is the core of the recommendation pipeline.
- **Hybrid Engine:** A conditional wrapper that determines whether to use CF, CB, or both based on user profile and search intent.
- **Evaluation Layer:** Generates prediction accuracy metrics and visual reports.
- **User Interface:** A prototype interface within Google Colab allowing user ID input and optional search terms.

## **Data Flow Diagram (DFD)**

### **Level 0:**

- User provides input → System fetches user/product data → Recommendations generated → Output displayed to user

### **Level 1:**

- Data collected from CSVs → Cleaned and merged → Split into training/testing sets → Passed to models → Model generates top N items

## **Use Case Diagram**

- **Actor:** End User
- **Use Cases:**
  - View recommended products
  - Search for similar products
  - Rate previously purchased products

## **UML Class Diagram**

- Classes: User, Product, Review, Rating, RecommenderEngine
- Methods: get\_top\_n(), predict\_rating(), calculate\_similarity(), train\_model()

## **Challenges Addressed in Project Design**

- Cold Start Problem: Handled using content-based recommendations based on product metadata.
- Sparsity of Ratings: Matrix factorization algorithms (like SVD) reduce dimensions and infer hidden patterns.
- High Dimensionality in Text: Applied TF-IDF with min\_df/max\_df filtering to control vector size.
- Data Integration: Combined data from over 10 separate CSV files ensuring referential integrity across IDs.

## **Benefits of the Proposed System**

- Scalable and adaptable to different product catalogs
- Personalized recommendations that evolve with user behavior
- Versatile architecture that allows adding new models or switching logic

## **CHAPTER 4: IMPLEMENTATION**

### **TOOLS AND TECHNOLOGIES USED**

The entire project was executed using Python as the core programming language. The team primarily worked on Google Colab, which provided a powerful cloud-based platform with free GPU access and seamless integration with Google Drive. This setup allowed for real-time collaboration and simplified version control among team members. Key libraries included Pandas, NumPy, Matplotlib, Seaborn, scikit-learn, and Surprise.

We used Pandas for data manipulation, NumPy for numerical operations, and Matplotlib and Seaborn for data visualization. For machine learning operations like TF-IDF vectorization, cosine similarity, and model evaluation, we used scikit-learn. For collaborative filtering, the Surprise library helped us train SVD, NMF, and KNNBasic models.

### **METHODOLOGY**

Our methodology was structured into multiple phases to ensure organized execution from data acquisition to model deployment:

1. **Data Integration:** We began by combining several raw CSV datasets, including customer information, order data, payment records, product metadata, and reviews. These were merged using unique identifiers such as `customer_id` and `order_id`.
2. **Data Cleaning & Preprocessing:** Irrelevant fields were removed, null values were handled appropriately, and text-based fields were cleaned to remove special characters. This created a robust and usable master dataset.
3. **Collaborative Filtering (CF):** We trained multiple collaborative filtering models (SVD, NMF, KNNBasic) using the Surprise library. The user-product-rating matrix was created, and each model was evaluated based on RMSE and MAE.
4. **Content-Based Filtering (CB):** We used TF-IDF vectorization on combined text fields (product name, brand, category, tags) and calculated cosine similarity to find semantically similar products.
5. **Hybrid System Development:** A logical condition was created to switch between CF and CB models based on user rating history. If a search keyword was entered, both CB and CF scores were weighted to generate final recommendations.
6. **UI and Output Handling:** A simple input method was created within the notebook to accept user IDs and keywords. The top-N recommended products were displayed based on predicted rating or similarity score.

### **DATA PREPROCESSING**

Preprocessing was a crucial part of this project. Initially, the raw datasets were inconsistent and contained numerous null values, irrelevant columns, and disjointed

tables. We merged multiple datasets such as customer info, order history, payments, reviews, and product details. This was followed by the removal of redundant fields like timestamps and zip codes.

We ensured that all missing values were handled either through imputation or by dropping rows that lacked significant information. Text data such as product descriptions and tags were cleaned by removing special characters and stop words. This cleaning process made the data usable for both collaborative and content-based recommendation models.

## **BUILDING THE COLLABORATIVE FILTERING MODEL**

We used the Surprise library to implement collaborative filtering. After formatting the dataset to contain `user_id`, `product_id`, and `rating`, we split it into training and testing sets. We trained three models: SVD, NMF, and KNNBasic. Each model was evaluated using RMSE and MAE.

The SVD model gave us the best results with the lowest RMSE and MAE, making it our primary recommendation engine for users with adequate rating history. Predictions were generated for unrated products for each user, and the top N highest-rated predictions were selected as recommendations.

## **BUILDING THE CONTENT-BASED FILTERING MODEL**

We created a custom TF-IDF-based model using scikit-learn. This model focused on textual features like product name, brand, category, and description. We generated TF-IDF vectors and calculated cosine similarity between products. This model worked especially well for users without prior history, solving the cold start problem efficiently.

Users could type in a search keyword (e.g., “smartphone”), and the system would return the most relevant products based on their textual metadata. This model was highly responsive to search context and gave meaningful results even for brand-new users.

## **HYBRID RECOMMENDATION STRATEGY**

To build a smart, user-adaptive system, we created a hybrid recommendation function. The logic worked as follows:

- If the user had more than 5 interactions, the system used collaborative filtering.
- If the user was new or had no history, content-based filtering was triggered.
- If a user also provided a keyword, both models contributed, with weights given to CF predictions.

This made the recommendation engine both intelligent and inclusive, handling all user types without compromising personalization quality.

## **USER INTERACTION AND INTERFACE**

We built a simple user interface in the Colab environment. Users could enter their unique user ID and an optional search term. The system would then return top 10 product recommendations along with predicted rating values. This made the engine interactive and simulated a realistic e-commerce experience for our evaluation.

This phase involved a lot of testing, fine-tuning, and experimentation, which contributed significantly to our understanding of model behaviors, performance metrics, and system design principles.

## MODULES / SCREENSHOTS

Our system was broken down into modular components for better manageability:

1. Data Loader Module – Responsible for importing, merging, and cleaning datasets.
2. Collaborative Filtering Engine – Trained using Surprise, this module handled predictions based on rating history.
3. Content-Based Engine – TF-IDF vectorizer with cosine similarity to calculate textual relevance.
4. Hybrid Controller – Logic handler to determine when to invoke CF, CB, or both.
5. User Interface Module – Accepted user inputs and displayed top recommendations.

## CODE SNIPPETS

Here are some sample snippets used in different parts of the project:

1. Loading and Preparing Dataset

```
import pandas as pd
orders = pd.read_csv('/content/orders.csv')
products = pd.read_csv('/content/products.csv')
data = pd.merge(orders, products, on='product_id')
```

2. TF-IDF Vectorization and Cosine Similarity

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
vectorizer = TfidfVectorizer()
tfidf_matrix = vectorizer.fit_transform(data['description'])
cosine_sim = cosine_similarity(tfidf_matrix)
```

3. Collaborative Filtering with SVD

```

from surprise import SVD, Dataset, Reader
reader = Reader(rating_scale=(1, 5))
data = Dataset.load_from_df(df[['user_id', 'product_id', 'rating'])
trainset = data.build_full_trainset()
model = SVD()
model.fit(trainset)

```

#### 4. Generating Hybrid Recommendations

```

def hybrid_recommend(user_id, keyword=None):
    if keyword:
        cb_result = content_based(keyword)
        cf_scores = [model.predict(user_id, pid) for pid in cb_result]
        return sorted(cf_scores, key=lambda x: x.est, reverse=True)
    else:
        return collaborative_only(user_id)

```

## CHAPTER 5: RESULTS AND DISCUSSION

### Collaborative Filtering Evaluation

We explored several collaborative filtering models, which primarily focus on users' past ratings and behaviors. Here's what we tested:

- SVD (Singular Value Decomposition): Captures underlying user-product preferences using matrix factorization.
- NMF (Non-negative Matrix Factorization): Similar to SVD but assumes all factors are non-negative, making it more interpretable.
- KNNBasic: A simpler algorithm that looks at nearest neighbors (similar users or products).
- BaselineOnly: Serves as a reference model using basic averages.

We measured the accuracy of these models using two industry-standard metrics:

- RMSE (Root Mean Squared Error)
- MAE (Mean Absolute Error)

Here's a snapshot of how each model performed:

Model	RMSE	MAE
SVD	0.94	0.74
NMF	0.98	0.79
BaselineOnly	0.96	0.76
KNNBasic	1.01	0.81

Among all models, SVD emerged as the most effective, delivering the lowest error margins. In simple terms, it was better at “guessing” how a user would rate a product they hadn't seen before. While KNNBasic was easier to understand, it didn't scale well with our dataset.

### Content-Based Filtering Evaluation

Next, we evaluated our content-based filtering system, which relies on the characteristics of the products themselves — such as name, description, brand, and tags. To analyze text data, we used TF-IDF vectorization, which helped us understand which words mattered most in distinguishing one product from another.

We tested different configurations:

- Trying 1-gram vs 2-gram word combinations
- Adjusting minimum and maximum document frequency thresholds

The best results came from using bigrams (1,2) and setting `min_df=0.005` and `max_df=0.99`. These settings struck the right balance: capturing enough product detail without overwhelming the system.

For new users (those who hadn't interacted with the system before), content-based recommendations became especially valuable. Based on visual review and informal testing, we observed that more than 80% of the recommendations were relevant and appealing.

## Hybrid Recommendation Model Evaluation

Here's where things got interesting. The hybrid engine combined both CF and CB methods intelligently. If a user had enough history (say, at least 5 ratings), the system used collaborative filtering. If not, or if the user typed a search query, it incorporated content-based logic.

### How We Measured Success:

To simulate a real online shopping experience, we randomly selected 500 users and generated Top-10 recommendations for each. Here's what we found:

- Hit Rate: 84% — meaning most users clicked or showed interest in at least one recommendation.
- Mean Reciprocal Rank (MRR): 0.41 — good engagement with higher-ranked items.
- Precision@10: 0.52
- Recall@10: 0.48

These results indicate that our hybrid model is both accurate and user-friendly.

## Visual Insights

To better interpret these results, we created several charts:

1. Most Popular Products: Showed the top 40 most interacted products.
2. Rating Distribution: Revealed a heavy skew toward 4-star and 5-star ratings.
3. User-Product Heatmap: Helped us visualize how ratings were distributed across users.
4. Top Categories: Electronics and fitness-related products were the most loved.
5. Rating vs Review Count: A scatter plot showed that higher review counts often correlated with slightly better ratings.



## Examples of Successful Recommendations

Let's look at a couple of real cases that illustrate how the system worked:

- **User ID: 198247**
  - Interests: Kitchen and Home Decor
  - Search: “wooden table”
  - Recommendations:
    - Modern Wooden Dining Table
    - Foldable Bamboo Study Desk
    - Wooden Storage Rack
- **User ID: 712305 (Cold Start)**
  - Search: “wireless earbuds”
  - Recommendations:
    - Noise Air Buds Pro 2
    - boAt Airdopes 141
    - JBL Tune 215TWS

In both cases, the recommendations felt tailored and aligned with user expectations, confirming that the hybrid system was working as intended.

## Challenges Faced

Like any real-world project, we ran into a few roadblocks:

1. Sparse Data
  - Some users rated very few products, making collaborative filtering difficult.
  - We addressed this by applying matrix factorization and ignoring low-activity users.
2. Cold Start Problem
  - When a brand-new user logs in, there's no history to rely on.
  - Solution: Lean heavily on content-based filtering for such users.
3. Computation Time

- TF-IDF matrices can get huge, especially with bigram features.
- We solved this by tuning the vectorizer and using sparse matrix storage.

#### 4. Duplicate Product Listings

- Slight differences in product names created confusion.
- We introduced basic text cleaning and filtering techniques.

### **Interpretation of Results**

The biggest takeaway? Flexibility matters. The hybrid engine didn't just stick to one logic — it adapted. This adaptability allowed us to serve users better, no matter their interaction history or preferences.

Also, allowing users to enter a search term gave them a sense of control over what they saw. It's not just about pushing products — it's about making the system feel helpful, like a smart assistant.

## **CHAPTER 6: CONCLUSION AND LEARNINGS**

### **PROJECT SUMMARY**

This internship project provided us with a hands-on opportunity to design, build, and test a complete recommendation system from scratch. Our goal was to create an intelligent e-commerce recommendation engine that could personalize suggestions based on either a user's past behavior or their keyword input. The system integrated collaborative filtering, content-based filtering, and a hybrid logic to support both frequent and new users. We successfully leveraged large, real-world datasets, implemented state-of-the-art algorithms, and developed a flexible architecture capable of adapting to multiple user scenarios.

The process involved integrating complex datasets from various CSV files, building robust preprocessing pipelines, training multiple recommendation models, evaluating their accuracy, and deploying them in a simplified UI environment using Google Colab. Our approach not only focused on technical correctness but also usability and scalability. This internship enabled us to experience the full lifecycle of a machine learning application.

### **LEARNINGS AND SKILLS GAINED**

Participating in this project significantly deepened our understanding of both theoretical and practical aspects of machine learning. We gained exposure to:

- **Data Engineering:** Combining and cleaning multiple datasets, feature extraction, handling missing values.
- **Exploratory Data Analysis:** Understanding patterns, distributions, and correlations through visualizations.
- **Model Training and Evaluation:** Training models like SVD, NMF, and KNNBasic and evaluating them with RMSE/MAE.
- **Text Mining and NLP:** Using TF-IDF and cosine similarity to measure textual product relevance.
- **Recommendation Systems:** Understanding how collaborative filtering, content-based filtering, and hybrid logic differ and when to apply each.
- **Python Programming:** Writing reusable, modular code and using libraries such as Pandas, NumPy, Matplotlib, scikit-learn, and Surprise.
- **Team Collaboration:** Dividing responsibilities, sharing codebases, and maintaining documentation.

We also improved soft skills such as problem-solving, debugging, version control, and communicating results effectively. Real-time testing, tuning models, and interpreting output for meaningful recommendations provided deep insight into data science workflows.

## **PERSONAL REFLECTIONS**

Working on this project made us realize that building a recommendation engine is more than just code—it requires understanding the user journey, business context, and data relationships. At each phase, we encountered challenges such as data sparsity, model overfitting, or handling missing reviews. Through trial and error, teamwork, and mentor support, we were able to solve them one by one.

Being part of a team of five allowed us to play to our strengths. Some members took the lead in model building, while others handled data engineering or visualization. This division of roles ensured that we not only completed the project on time but also understood the contributions of each system component.

The hybrid recommendation logic, in particular, was an eye-opener. It made us appreciate how real-world applications often rely on flexible systems that can adjust to user input dynamically. Watching our system generate different sets of recommendations for cold-start users vs. regular users was both technically satisfying and practically insightful.

## **APPLICATION AND FUTURE SCOPE**

This recommendation engine can be deployed on e-commerce platforms to assist customers in finding relevant products with greater accuracy. With further tuning and real-time data pipelines, it can be integrated into a scalable architecture.

We believe this system can be extended in the following ways:

- **User Behavior Analytics:** Track click-through rates and conversions to further refine recommendations.
- **Deep Learning Integration:** Use neural networks for text embeddings and collaborative filtering.
- **Real-Time Deployment:** Integrate APIs and a web dashboard for real-time usage.
- **Multilingual Search Support:** Enhance the TF-IDF system to support non-English inputs.

This project has built a strong foundation for our future careers in data science, AI, and software development. More importantly, it gave us the confidence to take on complex real-world problems with a structured and team-driven approach.

We conclude this internship experience with a profound appreciation for the power of recommendation systems and the real-world impact they create for users and businesses alike.